What is ethics anyway?

Gregory Johnson

1.

For something that pervades our lives, most people cannot give a very clear definition of ethics. That's easy to fix, however. Ethics is the study of morality. Of course, that only shifts the question to What is morality? Again, the definition is easily supplied. Actions—and maybe other things, but we'll focus on actions—are morally right or morally wrong. (Or neither, some actions are morally neutral.) If an action is morally right, then this means that we ought to do it. We might not, and we might not be punished if we don't, but we *should*. On the other hand, if the action is morally wrong, this means we should not do it. Still, when faced with a moral dilemma, knowing the definition of morality isn't really much help. Which action should I choose? And how do I know it's the right one?

One answer might be that religion and the associated religious texts tell us which actions are morally right and which ones are morally wrong. But using religion as a guide faces some problems. First, there's no consensus about which religion should be consulted. Christians will say Christianity, Hindus will say Hinduism, Jews will say Judaism, Muslims will say Islam, Buddhists will Buddhism, and on it goes; and that's not even getting into the sects that divide every religion. But even if we select one religious text—let's take the Christian Old Testament—upon examination, we find moral guidelines that don't seem quite right. There is slavery, the questionable treatment of women and children, and the command, given in *Exodus*, that "Six days shall work be done, but on the seventh day you shall have a holy Sabbath of solemn rest to the Lord; whoever does any work on it shall be put to death."

Making sense of the moral guidelines in the Old Testament can be set aside, however, because there is a deeper problem. One that was first articulated by the ancient Greek philosopher Plato in the fourth century B.C.E. Almost all of Plato's writings that we have today are dialogues, and in each the main character is Plato's slightly older contemporary Socrates. In a dialogue titled *Euthyphro*, Socrates and a priest named Euthyphro (who may or may not have been a real person) try to work out the definition of *piety*. About mid-way through the dialogue, Euthyphro proposes this definition "I would certainly say that the pious is what all the gods love, and the opposite, what all the gods hate, is the impious." Socrates, pressing for clarification, asks, "Is the pious being loved by the gods because it is pious, or is it pious because it is being loved by the gods?"

If we replace *the pious* with *morally correct actions* and *gods* with *God*, then Socrates is asking which one of these is correct:

(1) Morally correct actions are loved by God because they are morally correct.

(2) Morally correct actions are morally correct because they are loved by God.

If you can see the difference between these two, that's a good start. But let's pause for a moment to consider what's at stake here. The idea that the Old Testament (or any other religious text) tells us what is morally right and morally wrong is called the *divine command theory*. That's what's being invoked when someone says, "such-and-such is morally right because God (or the Bible) says that it is," and that's what Euthyphro has in mind with his definition. Plato is showing us that this idea can be understood in two ways. Today, when this dilemma is discussed, the two options are usually phrased this way:

(1b) God commands an action because it is moral.

(2b) An action is moral because God commands it.

The first option means that there are some morally correct actions, and God tells us to do those actions. But God didn't do anything to make those actions morally correct, and so they would still be morally correct even if God never existed. What's significant about this option is that actions are not right or wrong because God says that they are. Rather, there is some other reason why the actions are moral.

The second option states that God does make certain actions morally correct. By telling us to do certain things and not others, God thereby determines which actions are morally correct and which ones are morally wrong. Although at first glance, this looks like a powerful option, it turns out to be pretty unattractive. This is the scenario: imagine that all possible actions exist (killing, lying, telling the truth, helping the needy, and so forth), but none of them are morally right or morally wrong yet. Then God comes along and randomly (yes, randomly) chooses some that he is going to tell humans to do and some that he is going to tell humans not to do. If the Ten Commandments are to be believed, he settled on, among other things, not killing, not stealing, and not committing adultery. But he could just as well have well gone with kill, steal, and commit adultery. If he had, then killing, stealing, and adultery would all be morally correct. Of course, we might want to say that God wouldn't do that. God had reasons for choosing some actions to be moral and others to be immoral. But if God had reasons, then this option turns into the first one. With the second option, there aren't reasons; he just picked the actions.

Since it includes this element of randomness, the second option is not considered viable. But the first option tells us that the reason an action is morally correct is independent of what God says. So, morality has to have some justification other than the religious texts or God's word. The divine command theory falls short.

Still, although understanding why actions are morally right and morally wrong is an important part of doing ethics, we might think that we can still salvage something here by noting that, even if we don't have the reasons for why some actions are morally right and others are morally wrong, we can take comfort in the idea that God, with some justification or other, put various edicts in the Old Testament (or any other religious text) for us to find. We can simply proceed with his list. Unfortunately, even that isn't so simple, and the requirement that we kill people who work on the Sabbath isn't the only problem. In fact, let's set that aside, and take the commandment "thou shall not kill." Simple enough, but consider this situation described by the philosopher Bernard Williams:

> Jim finds himself in the central square of a small South American town. Tied up against the wall are a row of twenty Indians, most terrified, a few defiant, in front of them several armed men in uniform. A heavy man in a sweat-stained khaki shirt turns out to be the captain in charge and, after a good deal of questioning of Jim which establishes that he got there by accident while on a botanical expedition, explains that the Indians are a random group of the inhabitants who, after recent acts of protest against the government, are just about to be killed to remind other possible protesters of the advantages of not protesting. However, since Jim is an honored visitor from another land, the captain is happy to offer him a guest's privilege of killing one of the Indians himself. If Jim accepts, then as a special mark of the occasion, the other Indians will be let off. Of course, if Jim refuses, then there is no special occasion, and Pedro here will do what he was about to do when Jim arrived, and kill them all. Jim, with some desperate recollection of schoolboy fiction, wonders whether if he got hold of a gun, he could hold the captain, Pedro, and the rest of the soldiers to threat, but it is quite clear from the set-up that nothing

> of that kind is going to work: any attempt at that sort of thing will mean that all the Indians will be killed, and himself. The men against the wall, and the other villagers, understand the situation, and are obviously begging him to accept. What should he do?

Many people, although not all, have the strong intuition that Jim should kill one of the captives. We might recognize why Jim doesn't want to do this. We might even realize that we would have great difficulty doing it ourselves. But, nonetheless, it is what Jim *should* do. Killing in this instance is the morally correct action. Not taking the captain's offer would be morally wrong.

Now we are in the thick of it. We have a moral dilemma, and, perhaps, an idea of what the morally correct action is. But *why* is killing a person the right thing to do in this situation? Selecting an action is part of what we must do when faced with a moral dilemma, but it's only part of it. If we want to be able to say that we did the right thing, then we also have to justify our action. That is, we have to give the reasons why it was the correct action in that situation. The principle 'God commands an action because it is moral' does not provide those reasons, and so we turn to the first of the two main theories in ethics.

2. Utilitarianism

Utilitarianism is most commonly associated with the 19th century philosopher and politician John Stuart Mill. A statement of the theory is short and to the point: *the morally correct action is the one that produces the greatest amount of happiness for all involved, or all who will be affected*. Applying this theory just amounts to doing a calculation. First, we identify the different actions that can be taken and who will be affected by those actions. In the example above, Jim can either elect to kill one person or he can refuse. Jim's decision will affect, at least, the twenty captives, the other villagers, the captain, and Jim himself. Perhaps the circle can be expanded

further, but this will do. If Jim kills the one captive, then the other nineteen live. If Jim refuses, then Pedro will kill all twenty.

So, if Jim kills one person, there is much less loss of life. But that's not the only factor, there are also the villagers who, let's say, all have family and friends among the twenty captives. For some of them, Jim killing one person will cause a great deal of unhappiness. But those people will also be unhappy if Jim refuses and Pedro kills all twenty. Hence, if only one person is killed, the villagers, on the whole, will be much happier than if all twenty are executed. The captain seems as though he'll be equally content either way. That just leaves Jim.

The effect that killing or refusing to kill has on Jim counts, but it's no more or less important than the effect of these actions on anyone else. Maybe, if he kills the one person, Jim will eventually feel proud for doing something difficult and saving the lives that he could. Or maybe he'll be traumatized or feel guilty for the rest of his life. Maybe he'll be happier if he decides not to get involved. Or maybe he'll be wracked with guilt for the rest of his life if he refuses to kill. Whatever it might be is a factor, but it's not going to tip the balance, and so it really doesn't matter. If we add up all of the happiness and unhappiness that is created if Jim kills one captive and compare that to the amount of happiness and unhappiness created if Jim refuses, we find that, overall, Jim killing one person creates more happiness than if he refuses. Therefore, according to utilitarianism, killing the one person is the morally correct action.

Utilitarianism tells us what the morally correct action is, and it also tells us *why* that action is morally correct. But before thinking about the why, let's clarify one part of the theory: maximizing happiness. Sometimes simply adding up the number of people who are made happy and the number made unhappy will suffice, but that's not quite what the theory says. It's the amount of happiness or unhappiness that each action creates. Consider a different example.

> Imagine that you are a heart surgeon at a large hospital. One upcoming evening, you are not scheduled to work and so you make plans to meet six friends for dinner. They haven't seen you for many months, and you promise them that you will be there. The day of the dinner, however, every other heart surgeon at your hospital comes down with the flu and is sent home. Just as you are preparing to leave, a six-year-old girl who needs emergency heart surgery is admitted to the hospital. There is no one else who can perform this surgery, and so you must decide whether to break the promise that you made to your friends or to do the surgery. (For the sake of the example, let's say that the hospital won't compel you to do the surgery, and, for whatever reason, your friends won't ever know why you missed the dinner. Maybe it's a top-secret surgery.) The girl's family is limited to her and her parents, and they have no close friends.

If you perform the surgery, just the girl and her parents will be happy. Your six friends, meanwhile, will be unhappy. And to keep it simple, let's say that you will be equally happy either way.

Importantly, however, by performing the surgery and saving the girl's life, you will make her and her parents incredibly happy. Each one of your friend's unhappiness is far outweighed by the happiness of the girl or either one of her parents. Hence, even though fewer people are made happy than unhappy, performing the surgery will create more happiness than unhappiness. Conversely, if you join your friends for dinner, don't perform the surgery, and the girl dies, then the unhappiness that this creates, for her parents and briefly for the girl, will be enormous. Each of your friends' happiness will be tiny by comparison. Now, the girl's and her parents' unhappiness outweigh your six friends' happiness. Thus, even though performing the surgery will only make three people happy, according to utilitarianism, it is the morally correct action.

An issue that might occur to some people at this point is measurement. How do we measure and then add up happiness? There are methods for measuring happiness or, at least, for measuring the strength of people's preferences, although philosophers have noted certain problems with these methods. But often the specifics don't really matter. Everyone can see that more happiness is created if Jim kills one person than if he refrains and lets Pedro kill all twenty. Similarly, the heart surgery case is pretty straightforward. The only added twist is that if some people are made happy and others are made unhappy and, for each person, the strength the happiness or unhappiness varies, then we need to think a bit more about which sentiment outweighs the other to arrive at an answer.

Now, let's turn to the primary purpose of utilitarianism: providing a reason why Jim should kill the one captive. Happiness is valuable. In and of itself, it is a good thing. Unhappiness, pain, and suffering are not. Hence, according to utilitarianism, creating happiness and minimizing unhappiness, pain, and suffering are what give actions their moral worth. Utilitarianism takes that idea and gives us a formula for acting on it. Now, one need not accept this. There is no experiment or proof that tells us that actions that maximize happiness and minimize suffering are the only ones that have moral worth. But, nonetheless, utilitarianism supplies one justification for why some actions are morally correct and others are morally wrong. And, importantly, it seems like maximizing happiness is the right justification some of the time. The main problem for this theory is that, sometimes, utilitarianism tells us that an action is morally correct, when, to many people, it seems not to be. This is the classic example:

> A rape and murder, perhaps racially motivated, are committed in a town long beset by racial tension. You are the chief of police, and you've spent a lifetime working to make your community safe. Now this has happened, and an outbreak of violence, in which many people will probably be killed, looks likely. In turns out,

> however, that there is a homeless man in one of your jail cells. If you frame this man for the crime, there will be a quick trial, he will be found guilty, and the violence will be avoided. Besides this homeless man, no one but you and the real criminal—who will presumably remain silent—will know what you have done. What should you do?

In this case, the people affected by your decision are the homeless man, the real criminal, the people in the town, some of whom may die depending on your decision, and you, the chief of police. An outline of the calculation goes as follows. If you *don't frame* the innocent man, he won't be found guilty, but, probably, there will be an outbreak of violence and a number of people will die; there will also be longer lasting negative effects for the town. On the other hand, if you *frame* this man, he will go to prison and he may be put to death. He cares about whether this happens, but he may not have many friends or family who do. Everyone else in this community, meanwhile, will be relatively content and unharmed if he is quickly found guilty of the crimes. Either way, there's not going to be much, if any, happiness created. But, if you frame the homeless man, there will be much less unhappiness. Thus, according to utilitarianism, framing the innocent man is the morally correct action. It is what you, the chief of police, should do. The problem, though, is that, for many people, framing an innocent person seems wrong, not right, even given the consequences.

3. Kantian ethics

Our other major ethical theory is one that was developed by the 18th century philosopher Immanuel Kant. The theory is basically just a statement of what Kant called the *categorical imperative*. He formulated four versions of the categorical imperative, which he claimed were different ways of saying the same thing. The four are related, although it's not obvious that they are exactly equivalent. At any rate, this is the first

version: *act only in accordance with that maxim through which you can at the same time will that it become a universal law*.[1]

A maxim is a rule, and when using the categorical imperative, the first step is to identify the rule that you would be following if you performed some particular action. In the previous example, if you frame the homeless man, you are following this maxim: frame an innocent person. Next, with that maxim in hand, consider this: if you had the power to make this maxim a universal law, could you reasonably do it? The "reasonably" part is going to be important, but first, a universal law is, for instance, *the speed of light is 186,000 miles per second* or *energy equals mass times the speed of light squared* (i.e., $E = mc^2$). A little more informally, this is also a universal law: On earth or in any environment with a gravity similar to earth's, if you step off a ledge and nothing else interferes, you will drop down to the nearest surface. This isn't just a guideline, and it's not a law that you have the option of following or not. If you step off a ledge, you are going to drop. Similarly, (although only hypothetically) if you make *frame an innocent person* a universal law, every time you or anyone else has a chance to frame someone who is innocent, that's what will happen.

Kant thought that our ability to reason dictated morality. So the question is, if you could, would it be reasonable to make *frame an innocent person* a universal law? Kant's answer would be no, and that has nothing to do with whatever mayhem might be created or avoided by framing innocent people. Unlike utilitarianism, Kantian ethics puts no weight on the consequences of an action. Rather, it would not be reasonable to make *frame an innocent person* a universal law because someday you might be that

---

[1] The second version of the categorical imperative is equally important in contemporary ethics, but it will not be our focus. It states: "Act in such a way that you always treat humanity, whether in your own person or in the person of any other, never simply as a means, but always at the same time as an end."

innocent person who gets framed, and, naturally, you wouldn't want that to happen. Since you cannot wish for this maxim to be a universal law, it is morally wrong for you to frame the homeless man.

Conversely, both of the maxims: *search for criminals* and *do not frame innocent people* are ones that anyone could, reasonably, want to become universal laws. Hence, according to Kantian ethics, the morally correct actions are the ones that follow these maxims.

Again, we have (maybe) figured out which action is morally correct and which one is morally wrong. This also comes with a justification. All of this business about maxims and turning them into universal laws may seem odd, but it is a very clever formula for getting at an important idea. What you do in any particular situation should be what anyone else is also entitled to do. If no one else should do that action that you are considering, then you shouldn't either. Morality is holding yourself to the same rules that you want everyone else to follow.

But while the categorical imperative seems to get things right some of the time, just as with utilitarianism, there is a point at which it will conflict with many people's intuitions. This is famously illustrated by an objection to Kant's theory that was posed by the political writer Benjamin Constant in 1797. The example concerns a murderer "who has asked whether our friend, who is pursued by him, had taken refuge in our house." It seems obvious, at least to many people including Constant, that if a murderer is pursuing your friend—or pursuing anyone for that matter—you should direct the murderer away from, not toward, wherever his intended victim is hiding. But that's not what the categorical imperative tells us to do.

The maxim is *tell a lie*. Why this cannot, reasonably, be turned into a universal law is interesting. We might think that if it did become a universal law, then, just as in the example about framing an innocent person, it would put us at a disadvantage. Others would lie to us. That fits, but according to Kant, there is an even more fundamental reason why it

would not be reasonable to turn the maxim *tell a lie* into a universal law. If lying became a universal law, then it would no longer be possible to lie. No one would take anyone at their word, and so it would be impossible to effectively convey something that wasn't true. Making it a universal law would make it impossible to do what you are, right now, attempting to do. Because of the contradiction inherent in making it a universal law and, at the same time, acting on the maxim, we can't (reasonably) wish for that maxim to become a universal law. Hence, telling a lie is morally wrong. The maxim *tell the truth*, meanwhile, runs into no such problem, and so, according to the categorical imperative, telling the truth is always the morally correct action.

Nonetheless, lying to the murderer in this situation seems to be the only morally acceptable option. Even if we have a strong distaste for lying, one might think that we should overcome that distaste, and if we didn't, then we would be doing something morally wrong. But that's not what Kant thought. He stuck by the categorical imperative. Imagine, he said, you lie to the murder. But just at that moment, your friend, who had been hiding in your house, sneaks out and goes to where you, by lying, led the murder. Your well-intentioned lie only ensures that the murderer finds his victim.

Kant's response isn't too satisfying, however. If there is one thing that has to be true about ethics, it is that we make the best decisions that we can based on the facts that we have. Acting morally doesn't require us to be able to control the future, but it does require us to make reasonable inferences about what is going to happen and then act.

4. Rights

The one ethical concept that almost everyone has heard of and many people often invoke is *rights*. A natural response to the case of the sheriff thinking about framing an innocent man is to insist that the man has a right

not to be framed. Framing him infringes on his rights. As familiar as rights might be, though, stating the content of this theory, or perhaps theories, is not as simple as it is for utilitarianism or Kantian ethics. To even get started, we need to back up a bit. Earlier, I said that if an action is morally right, then it is something that you should do. You have an obligation—a moral obligation—to do it. Or, flipped around, if the action is morally wrong, then you have an obligation to refrain from doing it. Utilitarianism and Kantian ethics tell us which actions we are obligated to do and which we are obligated to refrain from doing. Rights are slightly different. If the homeless man has a right not to be framed, that doesn't impose an obligation on him. Having the right doesn't mean that *he* should keep himself from being framed. Rather, if he has the right not to be framed, then that imposes an obligation on others. It is because the man has this right that the police chief should not frame him, imprison him, and maybe let him be executed. In virtue of the man's right, it would be wrong for the police chief to act that way.

There are also different ways of talking about rights. There are *legal rights*—rights that people are granted, one way or another, by the state. There are also, maybe, *human* or *natural rights*. These are rights that people have simply because they are people. (Others besides human beings—some animals, for instance—may also have natural rights, but we'll have human beings in mind.) Thus, if the right to free speech is a human right, then all humans, everywhere and at all times, have this right. At some times in the past and in some places today, this right may not be recognized and trying to exercise the right is tricky or dangerous. But, still, everyone has it. If, on the other hand, the right to free speech is only a legal right, then in some places, people have that right, in other places they do not. In the United States, with some exceptions, we have it. In Turkmenistan, well, they're not so lucky.

One advantage of legal rights is that it is clear which ones people have. If it's written down and the government recognizes it, then the citizens have that right. But when rights are invoked with regard to whether an action is morally right or morally wrong, it's almost always human rights that we have in mind. It's a little bit of a mystery, though, how we confirm that people do, in fact, have human rights. And then, if they have them, which ones they have.

Let's assume that there are some human rights and look at another distinction: *negative rights* and *positive rights*. Earlier I said that rights impose an obligation on others. This obligation can work in two ways. On the one hand, it might just be an obligation not to interfere. If this is all that is required by the right, then it is a negative right. The standard example is the right to free speech. If I have that right, then everyone else has an obligation not to interfere with me as I speak. But that's all that others have to do. No one has to provide me with a podium or listen to what I'm saying. On the other hand, for some rights, having the right does impose an obligation on others to do something. Here the standard example is the right to an education. If children have this right, that doesn't only mean that no one should interfere with them as they get an education. It means that someone—maybe parents, maybe the community, maybe the state—has an obligation to provide them with an education. Or take the right to medical care, which is also a positive right. Despite all of the disagreements over this in the United States, we almost unanimously agree that people do have a right to some level of medical care. We don't want people to lie dying on park benches or to set their own broken bones at home. We agree that they are entitled to some amount of medical care, provided by someone.

With the distinction between positive and negative rights in mind, we can see that rights may give us an insight into what makes some actions morally correct and others morally incorrect. Take the claim that there are

human rights, but those rights are only negative rights. This is libertarianism, and, fundamentally, it is based on the idea that people should have as much freedom as possible. Our obligation is to let people do as they like, as long as they don't interfere with others. A competing idea is that there are, not only negative rights, but also some positive rights. This is based on the idea that there are some things—for instance, education and medical care—that are essential for a person to have a chance to flourish. Without those things, we would still, biologically, be human beings, but we would be mere animals or automatons. We wouldn't, fully, be persons—that is, rational, moral, autonomous agents. Hence, people have a right to those things, and someone or some group has an obligation to provide them.

The tension between the two positions is easy to see. Consider Judith Jarvis Thomson's example,

> In a small community, a child is suddenly struck with a deadly infection. There is a medicine that will fight the infection and save the child's life, but the only supply of it in this town is owned by and is in a locked box on the back porch of a woman who is away and cannot be contacted. There isn't time to obtain the medicine from somewhere else. What should be done?

The right to property is a negative right: as long as I acquire my property legally, everyone—including the state—has an obligation not to interfere with it. The right to some basic level of medical care is, as we said, a positive right: someone has an obligation to provide that care when it is needed. So here the woman's right to her property bumps up against the child's right to medical care. If there are only negative rights, then the child's right to medical care doesn't exist and it's morally wrong to take the medicine from the woman without her consent. If there are both positive and negative rights, then the child's right to medical care most likely exists, but there is still the question of whose right takes priority. In the end, all

we can say is that rights give us a means of addressing this dilemma, but if our starting point is just that there are some positive and some negative rights, then that alone will not produce a solution.

~ ≈ ~

Utilitarianism, Kantian ethics, and rights are three of the main players in the study of morality, but they are only part of it. There are variations of those theories and other theories altogether, including some that focus on the person we ought to be, not just the actions that we should or should not take. Virtue ethics, which was originally developed by the ancient Greek philosopher Aristotle, focuses on our character. According to Aristotle, our moral obligation is to cultivate virtues such as generosity, honesty, compassion, prudence, and courage. Once that's accomplished, hopefully, morally correct actions will follow. In a somewhat similar vein, in the second half of the twentieth century, a number of female philosophers developed the ethics of care, which orients morality around a more maternal perspective. Instead of rules that dictate what we should and should not do, the focus is on relationships, caring for and nurturing others, and making sacrifices for one's family or community.

5. Making moral decisions

On January 12, 2010, a 7.0 magnitude earthquake, followed almost immediately by 6.0 and 5.7 magnitude aftershocks, hit southern Haiti. The earthquake killed over 200,000 people and displaced more than a million. One of those affected was a thirty-eight-year-old Port-au-Prince resident named Nathalie LeBrun. Her house collapsed in the earthquake killing most of her extended family, but, as her luck would have it, she survived because she had checked into a hospital earlier that day. The earthquake didn't help, of course, but Nathalie's most significant medical problems were chronic conditions, severe heart failure and a related lung condition. In the week that followed the earthquake, Nathalie ended up at an

American-run field hospital where she was given oxygen, which made it possible for her to breathe normally and kept the oxygen level in her blood from falling too low.

The field hospital, however, had a very limited supply of bottled oxygen. The night after she arrived, the tank that Nathalie had been given ran out and she came close to dying. That outcome was averted when more bottled oxygen was found the next morning, and later, Nathalie was moved onto an oxygen concentrator—a device that could remove oxygen from the air and deliver it to a patient. But the situation was tenuous. The field hospital wasn't getting more bottled oxygen, and it also didn't have enough diesel, which was needed to run the generators that powered the oxygen concentrators.

After a couple of days, the field hospital's liaison officer, in consultation with the head doctor, decided that Nathalie would no longer be given oxygen, even though taking it from her meant that she would likely die. Other patients needed the oxygen during surgery, and they would then recover from their injuries. Nathalie, meanwhile, needed a constant supply of oxygen, and since she had a chronic condition, she might never be well enough to leave the hospital. The liaison officer was a captain and nurse practitioner named Patrick Kadilak. When Sheri Fink, who reported this story, asked him about Nathalie, he said,

> We're running out of oxygen. The country itself doesn't have oxygen. So, I have to make the decision, 'no, she can't have the oxygen; turn it off.' I have to look at the greater good that we can provide with the limited resources we have.

Since the American field hospital was no longer providing her with care, Nathalie had to be transferred to a Haitian hospital. The hospital to which she was being sent was unlikely to have oxygen, and so it was expected that she would probably die there.

When she arrived, she was in severe distress, but an American physician who was volunteering at the hospital, Dr. Paul Auerbach, improvised and treated her with the resources that he had: diuretics to remove the fluid from her lungs and a tank with a little bit of oxygen left in it. That stabilized her. Shortly thereafter, more fuel unexpectedly became available and Nathalie was put back on an oxygen concentrator. Several months later, with Sheri Fink's help, Nathalie traveled to the United States for surgery to correct her heart condition. Given its severity, however, that turned out not to be possible, and when she couldn't get a transplant, she died.

Treating Nathalie in the aftermath of the earthquake illustrates the tension between Kantian ethics and utilitarianism. We can reasonably wish for the maxim *save a life that is in danger* to become a universal law, and we cannot do the same for the maxim *don't save a life that is in danger*. Therefore, the categorical imperative tells us that the morally correct action is to do what we can to save Nathalie. We do the same for each patient who follows her, and, even if at some point we exhaust our resources, we've done what we could for each person as he or she was presented to us. Sheri Fink clearly leans toward Kantian ethics. When asked later about what we should take away from Nathalie's story. She answered,

> Let's not give up. The conclusion is let's not give up. It turned out there were options for this woman. It turns out that somebody was able to extend her life. Now you could very well argue that she should have died in that moment because look at all the resources that were spent. But I just feel like there was some value in her existence. There was so much value.

But at the same time, it was reasonable to believe—and likely true—that other patients would be saved if the oxygen was available for them. Given the crisis and the limited supply of oxygen, the utilitarian justification for withholding care from one person so that multiple other patients can

benefit (and in the end, benefit more than Nathalie would have) is straightforward. To maximize happiness, or at least to maximize positive outcomes, we withhold the oxygen from Nathalie and use it on the other patients who don't have chronic conditions. Now, that's the morally correct action.

So where does that leave us? These moral theories give us a procedure for determining which actions are morally correct and which are morally incorrect. It's not so easy, however, to figure out which theory is the right one. One option is to think long and hard about what, ultimately, gives actions their moral value. Maximizing happiness and minimizing suffering? Or holding ourselves to the same rules that we want everyone else to follow? Rights, virtue ethics, and the ethics of care provide us with more ideas and principles, and, for those who aren't faint of heart, there are many more ethical theories. If you come to the conclusion that one of these theories has gotten it right, you bite the bullet and stick with the theory no matter what.

Another option is to lean on our intuitions. In the example of Jim wandering into a South American village, killing the one captive seems to be the morally correct action. In the example of the sheriff in the town where a riot is imminent, not framing the innocent man seems to be the morally correct action. Or maybe not. People have different intuitions, which is one reason why we hope that an ethical theory will help us decide how to act. In the end, however, we probably don't have a perfect procedure for making moral decisions, and we haven't turned ethics into a science. But that shouldn't be taken to mean that whatever we feel like doing is thereby acceptable. The ethical theories give us insight into our moral decisions and the means to think about, discuss, and justify those decisions. Even if we selectively apply the theories, at least when we do, we understand why we think that the action we have taken is, in that situation,

the right one. And that is very different than just doing what we "feel" and refusing to give, or being unable to give, a clear justification for our action.

Recall that our original task was simply to understand what the terms *ethics* and *morality* mean. We've covered that much, and if it has turned out that understanding a little bit has only made things more complicated, then, well, we've learned that ethics isn't easy.