Week 2: Chapter 2, pp. 38 – 47 (Read the section on natural law theory. You don't have to read the section on Rawl's contract theory.)

If you want to read more about utilitarianism and Kantian ethics, see pp. 5 – 14 in this chapter: https://loighic.net/assets/pdf/intro/What_Is_Ethics_Anyway.pdf

Utilitarianism

Vaughn discusses two versions of utilitarianism: *act-utilitarianism* and *rule-utilitarianism*. The default theory is act-utilitarianism, and so anytime Vaughn (or anyone else) refers just to *utilitarianism*, they mean *act-utilitarianism*.

On p. 41 (and a little bit on p. 42), Vaughn applies act-utilitarianism and rule-utilitarianism to the case of a 10-year-old boy with cerebral palsy who is in constant pain and has the mental life of an infant. Notice that, for utilitarianism (and for Kantian ethics as we will see shortly), there are a series of steps that have to be taken to determine the morally correct action according to the theory. I'll go through the steps, but you should read what's in the chapter and if you want more explanation, you should look at "What is ethics anyway?"

Act-utilitarianism

For act-utilitarianism, we must, first, determine the possible actions.

(1) In this case, the possible actions are (*a*) keep Johnny in his current condition or (*b*) give Johnny a lethal injection (which will kill him).

Next, we must determine who will be affected by these actions and how they will be affected.

(2a) These people will be affected: Johnny, his parents, other family members, friends, and the medical staff. For this analysis, we'll assume that if the physician administers the lethal injection, he or she acts discreetly, and no one else becomes involved. (There is more to say about this assumption, but it's a plausible assumption with which to start.)

(2b) Vaughn explains how each person will be affected if Johnny is kept in his current condition, and how each will be affected if Johnny is given the lethal injection.

There are a few things to keep in mind about step (2b). First, although it is not necessary to be too precise, in principle, we need to identify the amount of happiness or suffering that each person will experience if action (*a*) is chosen and the amount that each will experience if action (*b*) is chosen.

Johnny and his parents will experience a lot of pain and suffering if he is kept in his current state. Friends, other family, and medical staff will also be unhappy if Johnny is kept in his current state, but their suffering will, presumably, be less. If we use a scale from -100 (maximum suffering) to +100 (maximum happiness) with zero being neutral, then we might get something like the values in the table right below. Although their suffering is different, if he is kept in his current condition, Johnny and his parents will each be at or near -100. Every other family member will be around -50, and so forth. On the far right is the total for that that row, and at the bottom of the last column is the "total" suffering that will be caused by keeping Johnny in his current state.

	number		
Johnny	1	-100	-100
his parents	2	-100	-200
other family	15	-50	-750
his parents' close friends	10	-30	-300
medical staff	6	-30	-180
total			**-1530**

Alternatively, if Johnny is given a lethal injection, then his suffering ends. His parents will, as Vaughn says, probably experience relief, as would the others who are involved. For this analysis, we will count relief as a positive experience—not happiness exactly, but going in that direction. So, we can make another table. Johnny is now gone, and so we don't give him a value. I assigned a lower value to his parents than to everyone else with the thought that, although his parents are relieved (and on the positive side of zero), they aren't really very happy. Everyone else, I imagine, would be able to be happier that Johnny is no longer suffering.

	number		
Johnny	1	-	-
his parents	2	12	24
other family	15	20	300
his parents' close friends	10	20	200
medical staff	6	20	120
total			**+644**

Different people would come up with different numbers, but no matter what they might be, we can see that keeping Johnny in his current state will create a lot of unhappiness and suffering (i.e., a large negative number) while giving him a lethal injection will create some amount of happiness (i.e., a positive number). Hence, according to act-utilitarianism, the morally correct action is giving Johnny the lethal injection.

In fact, we can see this even without making a table and assigning numbers to each person's experience. The table is sometimes useful, however, to remind us that we need to (1) identify each individual who will be affected, (2) think about how they will be affected, and (3) do something like a calculation.

Rule-utilitarianism

The analysis for rule-utilitarianism is similar, but with some important differences. Instead of considering the two actions that might be taken, we need to consider two rules.

(1) One of the rules is the rule that we would be following if we keep Johnny in his current state: (*a*) "Do not kill seriously impaired children, regardless of their suffering or the wishes of their parents."

The other rule is the one that we would be following if we give Johnny the lethal injection: (*b*) "Killing seriously impaired children is permissible if they are suffering severely and improvement is hopeless."

For the next step, we aren't focused on Johnny, his parents, or their particular situation at all. Instead, we need to consider the consequences if rule (*a*) is in effect (everywhere, or at least in the entire country), and the consequences if rule (*b*) is in effect. We're not going to identify individuals or count up amounts of happiness or unhappiness. Rather, we are just going to determine the likely consequences of having each rule in effect and whether those consequences would be good (i.e., creating happiness and relieving suffering) or bad (creating unhappiness and suffering).

(2) You can read the analysis that Vaughn has at the bottom of col. 2 on p. 41 in col. 1 on p. 42.

According to this theory, the morally correct *action* is the one that falls under the rule that would create the most happiness. So, for this case, according to rule-utilitarianism, the morally correct action is *not* giving Johnny the lethal injection.

One thing that you might notice about rule-utilitarianism (and this also applies to the first version of the categorical imperative in Kantian ethics) is that a lot depends on how we characterize the rules that cover that actions being considered. If we changed rule (*b*) slightly to include the reference to the parents' wishes, we would have this:

(b) Killing seriously impaired children is permissible if they are suffering severely, improvement is hopeless, and it is the wish of their (loving and devoted) parents.

Now, if we considered the consequences, we might find that this one has better consequences than rule (*a*).

Kantian ethics

Kantian ethics is a little more complicated, and it is a good idea to also read pp. 11 – 14 in "What is ethics anyway?" (https://loighic.net/assets/pdf/intro/What_Is_Ethics_Anyway.pdf)

I'm not going to repeat everything that is in "What is ethics anyway?" But here is how we would use the first version of the categorical imperative (top of col. 1 on p. 43) to determine the morally correct action in the Johnny case.

(1) Just like we did for rule-utilitarianism, we have to begin by determining the rule that we would be following if we did each action being considered. I'm going to change the rules (i.e., the maxims), however. These two will be better for illustrating how the first version of the categorical imperative works, and arguably, they fit better with how the categorical imperative is typically interpreted and applied.

(*a*) Do not kill impaired children.

(*b*) Kill impaired children.

(2) Next, we need to imagine being able to turn each of these rules into a law of nature. (See pp. 11 – 12 in "What is ethics anyway?" for what this means.)

(3) Next, we ask ourselves, if we could *reasonably* want to turn either rule into a law of nature. The *reasonably* part is important, and there are just two ways in which it would be unreasonable to want to turn the rule into a law of nature.

(i) If it would be bad for us if the rule became a law of nature, then it is unreasonable to want to make it a law of nature, or

(ii) If having the rule as a law of nature would make it impossible to do the action that we are considering, then it is unreasonable to want to make it a law of nature.

If it would be unreasonable to want to turn rule (*a*) or rule (*b*) into a law of nature for either reason, then the action that falls under that rule is morally wrong—according to Kantian ethics.

In this case, (ii) doesn't apply. If (*a*) became a law of nature, we could keep Johnny in his present state. And if (*b*) became a law of nature, we could give Johnny the lethal injection.

(i) does apply, although there could be some debate about how the analysis would go from here. But let's say that you think that it would be bad for you if (*b*) became a law of nature. You think this because if it became a law of nature, then someday your own child, or one of your nieces or nephews, or a friend's child could be given a lethal injection against your wishes. On the other hand, rule (*a*) runs into no such problem. If (*a*) became a law of nature, then life would go on the way that we mostly want it to go on—Johnny's unfortunate case aside.

The doctrine of double effect

Natural law theory is important in ethics, although for the most part, you get the same results with Kantian ethics (i.e., the same actions will be morally right and the same ones will be morally wrong), and Kantian ethics is generally considered the stronger theory. The doctrine of double effect is important, however, and it often comes up in bioethics, even when the details of the natural law theory are mostly skipped over.

The idea with this doctrine is that, in some cases, an action that initially seems to be morally wrong (according to the natural law theory) is, actually, morally acceptable as long as the *intention* of the person doing action is morally acceptable. The content of the doctrine is four tests, each one of which must be passed for the action to be morally acceptable.

The doctrine of double effect sometimes comes up in bioethics when there is a question about killing a severely suffering patient. Vaughn goes through an example of this on pp. 46 – 47. The two actions are very similar (or maybe even the same), but if it is interpreted one way, it is morally wrong (according to the natural law theory). If the action is interpreted in another way, it is morally acceptable. The two actions and how they are evaluated for the four tests are listed in the table.

	Giving the woman a lethal injection	**Giving the woman a large injection of morphine (which it is known will end her life)**
1.	This action—taking a life—is not morally permissible. This fails test 1.	This action—relieving her suffering—is morally permissible. This passes test 1.
2.	Ending the woman's life (a bad effect) in order to end her suffering (a good effect) fails test 2.	Relieving the woman's pain and suffering (a good effect) with the only intention of doing just that (there is no other goal) passes test 2.
3.	The woman's death (a bad effect) was intended. This fails test 3.	Her death from the injection was *foreseen*, but it wasn't *intended*. This passes test 3.
4.	Causing the death of an innocent person is a more significant bad effect than is relieving someone's pain and suffering. This fails test 4.	The good effect (relieving her suffering) is important enough to outweigh the bad effect that does occur (hastening her death). This passes test 4.