CHAPTER 1

Moral Reasoning in Bioethics

Any serious and rewarding exploration of bioethics is bound to be a challenging journey. What makes the trip worthwhile? As you might expect, this entire text is a long answer to that question. You therefore may not fully appreciate the trek until you have already hiked far along the trail. The short answer comes in three parts.

First, bioethics—like ethics, its parent discipline—is about morality, and morality is about life. Morality is part of the unavoidable, bittersweet drama of being persons who think and feel and choose. **Morality** concerns beliefs regarding morally right and wrong actions and morally good and bad persons or character. Whether we like it or not, we seem confronted continually with the necessity to deliberate about right and wrong, to judge someone morally good or bad, to agree or disagree with the moral pronouncements of others, to accept or reject the moral outlook of our culture or community, and even to doubt or affirm the existence or nature of moral concepts themselves. Moral issues are thus inescapable—including (or especially) those that are the focus of bioethics. In the twenty-first century, few can remain entirely untouched by the pressing moral questions of fair distribution of health care resources, abortion and infanticide, euthanasia and assisted suicide, exploitative research on children and populations in developing countries, human cloning and genetic engineering, assisted reproduction and surrogate parenting, prevention and treatment of HIV/AIDS, the confidentiality and consent of patients, the refusal of medical treatment on religious grounds, experimentation on human embryos and fetuses, and the just allocation of scarce life-saving organs.

Second, it would be difficult to imagine moral issues more important—more closely gathered around the line between life and death, health and illness, pain and relief, hope and despair—than those addressed by bioethics. Whatever our view of these questions, there is little doubt that they matter immensely. Whatever answers we give will surely have weight, however they fall.

Third, as a systematic study of such questions, bioethics holds out the possibility of answers. The answers may or may not be to our liking; they may confirm or confute our preconceived notions; they may take us far or not far enough. But, as the following pages will show, the trail has more light than shadow—and thinking critically and carefully about the problems can help us see our way forward.

ETHICS AND BIOETHICS

Morality is about people's moral judgments, principles, rules, standards, and theories—all of which help direct conduct, mark out moral practices, and provide the yardsticks for measuring moral worth. We use *morality* to refer generally to these aspects of our lives (as in "Morality is essential") or more specifically to the beliefs or practices of particular groups or persons (as in "American morality" or "Kant's morality"). *Moral*, of course, pertains to morality as just defined, though it is also sometimes employed as a synonym for right or good, just as *immoral* is often meant to be equivalent to wrong or bad. *Ethics*, as used in this text, is not synonymous with *morality*. **Ethics** is the study of morality using the tools and methods of

philosophy. Philosophy is a discipline that systematically examines life's big questions through critical reasoning, logical argument, and careful reflection. Thus, ethics—also known as *moral philosophy*—is a reasoned way of delving into the meaning and import of moral concepts and issues and of evaluating the merits of moral judgments and standards. (As with *morality* and *moral*, we may use *ethics* to say such things as "Kant's ethics" or may use *ethical* or *unethical* to mean right or wrong, good or bad.) Ethics seeks to know whether an action is right or wrong, what moral standards should guide our conduct, whether moral principles can be justified, what moral virtues are worth cultivating and why, what ultimate ends people should pursue in life, whether there are good reasons for accepting a particular moral theory, and what the meaning is of such notions as *right*, *wrong*, *good*, and *bad*. Whenever we try to reason carefully about such things, we enter the realm of ethics: We *do* ethics.

Science offers another way to study morality, and we must carefully distinguish this approach from that of moral philosophy. **Descriptive ethics** is the study of morality using the methodology of science. Its purpose is to investigate the *empirical* facts of morality—the actual beliefs, behaviors, and practices that constitute people's moral experience. Those who carry out these inquiries (usually anthropologists, sociologists, historians, and psychologists) want to know, among other things, what moral beliefs a person or group has, what caused the subjects to have them, and how the beliefs influence behavior or social interaction. Very generally, the difference between ethics and descriptive ethics is this: In ethics we ask, as Socrates did, *How ought we to live?* In descriptive ethics we ask, *How do we in fact live?*

Ethics is a big subject, so we should not be surprised that it has three main branches, each dealing with more or less separate but related sets of ethical questions. **Normative ethics** is the search for, and justification of, moral standards, or norms. Most often the standards are moral principles, rules, virtues, and theories, and the lofty aim of this branch is to establish rationally some or all of these as proper guides for our actions and judgments. In normative ethics, we ask questions like these: What moral principles, if any, should inform our moral judgments? What role should virtues play in our lives? Is the principle of autonomy justified? Are there any exceptions to the moral principle of "do not kill"? How should we resolve conflicts between moral norms? Is contractarianism a good moral theory? Is utilitarianism a better theory?

A branch that deals with much deeper ethical issues is metaethics. **Metaethics** is the study of the meaning and justification of basic moral beliefs. In normative ethics we might ask whether an action is right or whether a person is good, but in metaethics we would more likely ask *what it means* for an action to be right or for a person to be good. For example, does *right* mean *has the best consequences*, or *produces the most happiness*, or *commanded by God*? It is the business of metaethics to explore these and other equally fundamental questions: What, if anything, is the difference between moral and nonmoral beliefs? Are there such things as moral facts? If so, what sort of things are they, and how can they be known? Can moral statements be true or false—or are they just expressions of emotions or attitudes without any truth value? Can moral norms be justified or proven?

The third main branch is **applied ethics**, the use of moral norms and concepts to resolve practical moral issues. Here, the usual challenge is to employ moral principles, theories, arguments, or analyses to try to answer moral questions that confront people every day. Many such questions relate to a particular professional field such as law, business, or journalism, so we have specialized subfields of applied ethics like legal ethics, business ethics, and journalistic ethics. Probably the largest and most energetic subfield is bioethics.

Bioethics is applied ethics focused on health care, medical science, and medical technology. (*Biomedical ethics* is often used as a synonym,

and *medical ethics* is a related but narrower term used most often to refer to ethical problems in medical practice.) Ranging far and wide, bioethics seeks answers to a vast array of tough ethical questions: Is abortion ever morally permissible? Is a woman justified in having an abortion if prenatal genetic testing reveals that her fetus has a developmental defect? Should people be allowed to select embryos by the embryos' sex or other genetic characteristics? Should human embryos be used in medical research? Should human cloning be prohibited? Should physicians, nurses, physicians' assistants, and other health care professionals always be truthful with patients whatever the consequences? Should severely impaired newborns be given life-prolonging treatment or be allowed to die? Should people in persistent vegetative states be removed from life support? Should physicians help terminally ill patients commit suicide? Is it morally right to conduct medical research on patients without their consent if the research would save lives? Should human stem-cell research be banned? How should we decide who gets life-saving organ transplants when usable organs are scarce and many patients who do not get transplants will die? Should animals be used in biomedical research?

The ethical and technical scope of bioethics is wide. Bioethical questions and deliberations now fall to nonexpert and expert alike—to patients, families, and others as well as to philosophers, health care professionals, lawyers, judges, scientists, clergy, and public policy specialists. Though the heart of bioethics is moral philosophy, fully informed bioethics cannot be done without a good understanding of the relevant nonmoral facts and issues, especially the medical, scientific, technological, and legal ones.

ETHICS AND THE MORAL LIFE

Morality, then, is a normative, or evaluative, enterprise. It concerns moral norms or standards that help us decide the rightness of actions, judge the goodness of persons or character, and prescribe the form of moral conduct. There are, of course, other sorts of norms we apply in life—*nonmoral* norms. Aesthetic norms help us make value judgments about art; norms of etiquette about polite social behavior; grammatical norms about correct use of language; prudential norms about what is in one's interests; and legal norms about lawful and unlawful acts. But moral norms differ from these nonmoral kinds. Some of the features they are thought to possess include the following.

Normative Dominance. In our moral practice, moral norms are presumed to dominate other kinds of norms, to take precedence over them. Philosophers call this characteristic of moral norms *overridingness* because moral considerations so often seem to override other factors. A maxim of prudence, for example, may suggest that you should steal if you can avoid getting caught, but a moral prohibition against stealing would overrule such a principle. An aesthetic (or pragmatic) norm implying that homeless people should be thrown in jail for blocking the view of a beautiful public mural would have to yield to moral principles demanding more humane treatment of the homeless. A law mandating brutal actions against a minority group would conflict with moral principles of justice and would therefore be deemed illegitimate. We usually think that immoral laws are defective, that they need to be changed, or that, in rare cases, they should be defied through acts of civil disobedience.

Universality. Moral norms (but not *exclusively* moral norms) have universality: Moral principles or judgments apply in all relevantly similar situations. If it is wrong for you to tell a lie in a particular circumstance, then it is wrong for everyone in relevantly similar circumstances to tell a lie. Logic demands this sort of consistency. It makes no sense to say that Maria's doing action A in circumstances C is morally wrong, but John's doing A in circumstances relevantly similar to C is morally right. Universality,

however, is not unique to moral norms; it's a characteristic of all normative spheres.

Impartiality. Implicit in moral norms is the notion of impartiality—the idea that everyone should be considered equal, that everyone's interests should count the same. From the perspective of morality, no person is any better than any other. Everyone should be treated the same unless there is a morally relevant difference between persons. We probably would be completely baffled if someone seriously said something like "murder is wrong . . . except when committed by myself," when there was no morally relevant difference between that person and the rest of the world. If we took such a statement seriously at all, we would likely not only reject it but also would not even consider it a bona fide moral statement.

The requirement of moral impartiality prohibits discrimination against people merely because they are different—different in ways that are not morally relevant. Two people can be different in many ways: skin color, weight, gender, income, age, occupation, and so forth. But these are not differences relevant to the way they should be treated as persons. However, if there are morally relevant differences between people, then we may have good reasons to treat them differently, and this treatment would not be a violation of impartiality. This is how philosopher James Rachels explains the point:

> The requirement of impartiality, then, is at bottom nothing more than a proscription against arbitrariness in dealing with people. It is a rule that forbids us from treating one person differently from another *when there is no good reason to do so.* But if this explains what is wrong with racism, it also explains why, in some special kinds of cases, it is not racist to treat people differently. Suppose a film director was making a movie about the life of Martin Luther King, Jr. He would have a perfectly good reason for ruling out Tom Cruise for the starring role. Obviously, such casting would make no sense. Because there would be a good reason for it, the director's "discrimination" would not be arbitrary and so would not be open to criticism.[1]

Reasonableness. To participate in morality—to engage in the essential, unavoidable practices of the moral life—is to do moral reasoning. If our moral judgments are to have any weight at all, if they are to be anything more than mere personal taste or knee-jerk emotional response, they must be backed by the best of reasons. They must be the result of careful reflection in which we arrive at good reasons for accepting them, reasons that could be acknowledged as such by any other reasoning persons.

Both logic and our commonsense moral experience demand that the thorough sifting of reasons constitutes the main work of our moral deliberations—regardless of our particular moral outlook or theory. We would think it odd, perhaps even perverse, if someone asserted that physician-assisted suicide is always morally wrong—and then said she has no reasons at all for believing such a judgment but *just does*. Whatever our views on physician-assisted suicide, we would be justified in ignoring her judgment, for we would have no way to distinguish it from personal whim or wishful thinking. Likewise she herself (if she genuinely had no good reasons for her assertion) would be in the same boat, adrift with a firm opinion moored to nothing solid.

Our feelings, of course, are also part of our moral experience. When we ponder a moral issue we care about (abortion, for example), we may feel anger, sadness, disgust, fear, irritation, or sympathy. Such strong emotions are normal and often useful, helping us empathize with others, deepening our understanding of human suffering, and sharpening our insight into the consequences of our moral decisions. But our feelings can mislead us by reflecting not moral truth but our own psychological needs, our own personal or cultural biases, or our concern for personal advantage. Throughout history, some people's feelings led them to conclude that women should be burned for witchcraft, that

IN DEPTH
MORALITY AND THE LAW

Some people confuse morality with the law, or identify the one with the other, but the two are distinct though they may often coincide. Laws are norms enacted or enforced by the state to protect or promote the public good. They specify which actions are *legally* right or wrong. But these same actions can also be judged *morally* right or wrong, and these two kinds of judgments will not necessarily agree. Lying to a friend about a personal matter, deliberately trying to destroy yourself through reckless living, or failing to save a drowning child (when you easily could have) may be immoral—but not illegal. Racial bias, discrimination based on gender or sexual orientation, slavery, spousal rape, and unequal treatment of minority groups are immoral—but, depending on the society, they may not be illegal.

Much of the time, however, morality and the law overlap. Often what is immoral also turns out to be illegal. This is usually the case when immoral actions cause substantial harm to others, whether physical or economic. Thus, murder and embezzlement are both immoral and illegal, backed by social disapproval and severe sanctions imposed by law. Controversy often arises when an action is not obviously or seriously harmful but is considered immoral by some who want the practice prohibited by law. The contentious notion at work is that something may be made illegal solely on the grounds that it is immoral, regardless of any physical or economic harm involved. This view of the law is known as *legal moralism*, and it sometimes underlies debates about the legalization of abortion, euthanasia, reproductive technology, contraception, and other practices.

Many issues in bioethics have both a moral and legal dimension, and it is important not to confuse the two. Sometimes the question at hand is a moral one (whether, for example, euthanasia is ever morally permissible); whether a practice should be legal or illegal then is beside the point. Sometimes the question is about legality. And sometimes the discussion concerns both. A person may consider physician-assisted suicide morally acceptable but argue that it should nevertheless be illegal because allowing the practice to become widespread would harm both patients and the medical profession.

whole races should be exterminated, that Black men should be lynched, and that adherents of a different religion were evil. Critical reasoning can help restrain such terrible impulses. It can help us put our feelings in proper perspective and achieve a measure of impartiality. Most of all, it can guide us to moral judgments that are trustworthy because they are supported by the best of reasons.

The moral life, then, is about grappling with a distinctive class of norms marked by normative dominance, universality, impartiality, and reasonableness. As we saw earlier, these norms can include moral principles, rules, theories, and judgments. We should notice that we commonly apply these norms to two distinct spheres of our moral experience—to both moral *obligations* and moral *values*.

Moral obligations concern our duty, what we are obligated to do. That is, obligations are about conduct, how we ought or ought not to behave. In this sphere, we talk primarily about *actions*. We may look to moral principles or rules to guide our actions, or study a moral theory that purports to explain right actions, or make judgments about right or wrong actions.

Moral values, however, generally concern those things that we judge to be morally good, bad, praiseworthy, or blameworthy. Normally we use such words to describe persons (as in "He is a good person" or "She is to blame for hurting them"), their character ("He is virtuous"; "She is honest"), or their motives ("She did wrong but did not mean to"). Note that we also attribute *nonmoral* value to things. If we say that a book or bicycle or vacation is good, we mean good in

a nonmoral sense. Such things in themselves cannot have *moral* value.

Strictly speaking, only actions are morally *right* or *wrong*, but persons are morally *good* or *bad* (or some degree of goodness or badness). With this distinction we can acknowledge a simple fact of the moral life: A good person can do something wrong, and a bad person can do something right. A Gandhi can tell a lie, and a Hitler can save a drowning man.

In addition, we may judge an action right or wrong depending on the motive behind it. If John knocks a stranger down in the street to prevent her from being hit by a car, we would deem his action right (and might judge him a good person). But if he knocks her down because he dislikes the color of her skin, we would believe his action wrong (and likely think him evil).

The general meaning of *right* and *wrong* seems clear to just about everyone. But we should be careful to differentiate degrees of meaning in these moral terms. *Right* can mean either "obligatory" or "permissible." An obligatory action is one that would be wrong *not* to perform. We are obligated or required to do it. A permissible action is one that is permitted. It is not wrong to perform it. *Wrong* means "prohibited." A prohibited action is one that would be wrong to perform. We are obligated or required *not* to do it. A *supererogatory* action is one that is "above and beyond" our duty. It is praiseworthy—a good thing to do—but not required. Giving all your possessions to the poor is generally considered a supererogatory act.

MORAL PRINCIPLES IN BIOETHICS

As noted earlier, the main work of bioethics is trying to solve bioethical problems using the potent resources and methods of moral philosophy, which include, at a minimum, critical reasoning, logical argument, and conceptual analysis. Many, perhaps most, moral philosophers would be quick to point out that beyond these tools of reason we also have the considerable help of moral principles. (The same could be said about moral theories, which we explore in the next chapter.) Certainly to be useful, moral principles must be interpreted, often filled out with specifics, and balanced with other moral concerns. But both in everyday life and in bioethics, moral principles are widely thought to be indispensable to moral decision-making.

We can see appeals to moral principles in countless cases. Confronted by a pain-racked, terminally ill patient who demands to have his life ended, his physician refuses to comply, relying on the principle that "it is wrong to intentionally take a life." Another physician makes a different choice in similar circumstances, insisting that the relevant principle is "ending the suffering of a hopelessly ill patient is morally permissible." An infant is born anencephalic (without a brain); it will never have a conscious life and will die in a few days. The parents decide to donate the infant's organs to other children so they might live, which involves taking the organs right away before they deteriorate. A critic of the parents' decision argues that "it is unethical to kill in order to save." But someone else appeals to the principle "save as many children as possible."[2] In such ways moral principles help guide our actions and inform our judgments about right and wrong, good and evil.

As discussed in Chapter 2, moral principles are often drawn from a moral theory, which is a moral standard on the most general level. The principles are derived from or supported by the theory. Many times we simply appeal directly to a plausible moral principle without thinking much about its theoretical underpinnings.

Philosophers make a distinction between absolute and prima facie principles (or duties). An *absolute* principle applies without exceptions. An absolute principle that we should not lie demands that we never lie regardless of the circumstances or the consequences. In contrast, a *prima facie* principle applies in all cases unless an exception is warranted. Exceptions are justified when the principle conflicts with other principles and is thereby overridden. W. D. Ross is given credit for drawing this distinction in his

1930 book *The Right and the Good*.[3] It is essential to his account of ethics, which has a core of several moral principles or duties, any of which might come into conflict.

Physicians have a prima facie duty to be truthful to their patients as well as a prima facie duty to promote their welfare. But if these duties come in conflict—if, for example, telling a patient the truth about his condition would somehow result in his death—a physician might decide that the duty of truthfulness should yield to the weightier duty to do good for the patient.

Moral principles are many and varied, but in bioethics the following have traditionally been extremely influential and particularly relevant to the kinds of moral issues that arise in health care, medical research, and biotechnology. In fact, many—perhaps most—of the thorniest issues in bioethics arise from conflicts among these basic principles. In one formulation or another, each one has been integral to major moral theories, providing evidence that the principles capture something essential in our moral experience. The principles are (1) autonomy, (2) nonmaleficence, (3) beneficence, (4) utility, and (5) justice.[4]

Autonomy

Autonomy refers to a person's rational capacity for self-governance or self-determination—the ability to direct one's own life and choose for oneself. The principle of autonomy insists on full respect for autonomy. One way to express the principle is: *Autonomous persons should be allowed to exercise their capacity for self-determination*. According to one major ethical tradition, autonomous persons have intrinsic worth precisely because they have the power to make rational decisions and moral choices. They therefore must be treated with respect, which means not violating their autonomy by ignoring or thwarting their ability to choose their own paths and make their own judgments.

The principle of respect for autonomy places severe restraints on what can be done to an autonomous person. There are exceptions, but in general we are not permitted to violate people's autonomy just because we disagree with their decisions, or because society might benefit, or because the violation is for their own good. We cannot legitimately impair someone's autonomy without strong justification for doing so. Conducting medical experiments on patients without their consent, treating competent patients against their will, physically restraining or confining patients for no medical reason—such practices constitute obvious violations of personal autonomy.

Not all restrictions on autonomy, however, are of the physical kind. Autonomy involves the capacity to make personal choices, but choices cannot be considered entirely autonomous unless they are fully informed. When we make decisions in ignorance—without relevant information or blinded by misinformation—our autonomy is diminished just as surely as if someone physically manipulated us. If this is correct, then we have a plausible explanation of why lying is generally prohibited: Lying is wrong because it undermines personal autonomy. Enshrined in bioethics and in the law, then, is the precept of *informed consent*, which demands that patients be allowed to freely consent to or decline treatments and that they receive the information they need to make informed judgments about them.

In many ways, autonomy is a delicate thing, easily compromised and readily thwarted. Often a person's autonomy is severely undermined not by other people but by nature, nurture, or his or her own actions. Some drug addicts and alcoholics, people with serious psychiatric illness, and those with severe mental impairment are thought to have drastically diminished autonomy (or to be essentially nonautonomous). Bioethical questions then arise about what is permissible to do to them and who will represent their interests or make decisions regarding their care. Infants and children are also not fully autonomous, and the same sorts of questions are forced on parents, guardians, and health care workers.

Like all the other major principles discussed here, respect for autonomy is thought to be prima facie. It can sometimes be overridden by considerations that seem more important or compelling—considerations that philosophers and other thinkers have formulated as principles of autonomy restriction. The principles are articulated in various ways, are applied widely to all sorts of social and moral issues, and are themselves the subject of debate. Chief among these is the harm principle: a person's autonomy may be curtailed to prevent harm to others. To prevent people from being victimized by thieves and murderers, we have a justice system that prosecutes and imprisons the perpetrators. To discourage hospitals and health care workers from hurting patients through carelessness or fraud, laws and regulations limit what they can do to people in their care. To stop someone from spreading a deadly, contagious disease, health officials may quarantine him against his will.

Another principle of autonomy restriction is paternalism. **Paternalism** is the overriding of a person's actions or decision-making for her own good. Some cases of paternalism (sometimes called *weak paternalism*) seem permissible to many people—when, for example, seriously depressed or psychotic patients are temporarily restrained to prevent them from injuring or killing themselves. Other cases are more controversial. Researchers hoping to develop a life-saving treatment give an experimental drug to someone without his knowledge or consent. Or a physician tries to spare the feelings of a competent, terminally ill patient by telling her that she will eventually get better, even though she insists on being told the truth. The paternalism in such scenarios (known as *strong paternalism*) is usually thought to be morally objectionable. Many controversies in bioethics center on the morality of strong paternalism.

Nonmaleficence

The principle of **nonmaleficence** asks us not to intentionally or unintentionally inflict harm on others. In bioethics, nonmaleficence is the most widely recognized moral principle. Its aphoristic expression has been embraced by practitioners of medicine for centuries: "Above all, do no harm." A more precise formulation of the principle is: *We should not cause unnecessary injury or harm to those in our care*. In whatever form, nonmaleficence is the bedrock precept of countless codes of professional conduct, institutional regulations, and governmental rules and laws designed to protect the welfare of patients.

A health care professional violates this principle if he or she deliberately performs an action that harms or injures a patient. If a physician intentionally administers a drug that she knows will induce a heart attack in a patient, she obviously violates the principle—she clearly does something that is morally (and legally) wrong. But she also violates it if she injures a patient through recklessness, negligence, or inexcusable ignorance. She may not intend to hurt anyone, but she is guilty of the violation just the same.

Implicit in the principle of nonmaleficence is the notion that health professionals must exercise "due care." The possibility of causing some pain, suffering, or injury is inherent in the care and treatment of patients, so we cannot realistically expect health professionals never to harm anyone. But we do expect them to use due care—to act reasonably and responsibly to minimize the harm or the chances of causing harm. If a physician must cause patients some harm to effect a cure, we expect her to try to produce the least amount of harm possible to achieve the results. And even if her treatments cause no actual pain or injury in a particular instance, we expect her not to use treatments that have a higher chance of causing harm than necessary. By the lights of the nonmaleficence principle, subjecting patients to unnecessary risks is wrong even if no damage is done.

Beneficence

The principle of *beneficence* has seemed to many to constitute the very soul of morality—or very close to it. In its most general form, it says that *we should do good to others*. (*Benevolence* is

different, referring more to an attitude of goodwill toward others than to a principle of right action.) *Beneficence* enjoins us to advance the welfare of others and prevent or remove harm to them.

Beneficence demands that we do more than just avoid inflicting pain and suffering. It says that *we should actively promote the well-being of others and prevent or remove harm to them*. In bioethics, there is little doubt that physicians, nurses, researchers, and other professionals have such a duty. After all, helping others, promoting their good, is a large part of what these professionals are obliged to do.

But not everyone thinks that *we all* have a duty of active beneficence. Some argue that though there is a general (applicable to all) duty not to harm others, there is no general duty to help others. They say we are not obligated to aid the poor, feed the hungry, or tend to the sick. Such acts are not required, but are supererogatory, beyond the call of duty. Others contend that though we do not have a general duty of active beneficence, we are at least sometimes obligated to look to the welfare of people we care about most—such as our parents, children, spouses, and friends. In any case, it is clear that in certain professions—particularly medicine, law, and nursing—benefiting others is often not just supererogatory but obligatory and basic.

Utility

The principle of *utility* says that *we should produce the most favorable balance of good over bad (or benefit over harm) for all concerned.* The principle acknowledges that in the real world, we cannot always *just* benefit others or *just* avoid harming them. Often we cannot do good for people without also bringing them some harm, or we cannot help everyone who needs to be helped, or we cannot help some without also hurting or neglecting others. In such situations, the principle says, we should do what yields the best overall outcome—the maximum good and minimum evil, everyone considered. The utility principle, then, is a supplement to, not a substitute for, the principles of autonomy, beneficence, and justice.

In ethics this maxim comes into play in several ways. Most famously it is the defining precept of the moral theory known as utilitarianism (discussed in Chapter 2). But it is also a standalone moral principle applied everywhere in bioethics to help resolve the kind of dilemmas just mentioned. A physician, for example, must decide whether a treatment is right for a patient, and that decision often hinges on whether the possible benefits of the treatment outweigh its risks by an acceptable margin. Suppose a man's clogged artery can be successfully treated with open-heart surgery, a procedure that carries a considerable risk of injury and death. But imagine that the artery can also be successfully opened with a regimen of cholesterol-lowering drugs and a low-fat diet, both of which have a much lower chance of serious complications. The principle of utility seems to suggest that the latter course is best and that the former is morally impermissible.

The principle also plays a major role in the creation and evaluation of the health policies of institutions and society. In these large arenas, most people aspire to fulfill the requirements of beneficence and maleficence, but they recognize that perfect beneficence or maleficence is impossible: Trade-offs and compromises must be made, scarce resources must be allotted, help and harm must be balanced, life and death must be weighed—tasks almost always informed by the principle of utility.

Suppose, for example, we want to mandate the immunization of all schoolchildren to prevent the spread of deadly communicable diseases. The cost in time and money will be great, but such a program could save many lives. There is a down side, however: A small number of children—perhaps as many as 2 for every 400,000 immunizations—will die because of a rare allergic reaction to the vaccine. It is impossible to predict who will have such a reaction (and impossible to prevent it), but it is almost certain to occur in a few cases. If our goal is

social beneficence, what should we do? Children are likely to die whether we institute the program or not. Guided by the principle of utility (as well as other principles), we may decide to proceed with the program since many more lives would likely be saved by it than lost because of its implementation.

Again, suppose governmental health agencies have enough knowledge and resources to develop fully a cure for only one disease—either a rare heart disorder or a common form of skin cancer. Trying to split resources between these two is sure to prevent development of any cure at all. The heart disorder kills 200 adults each year; the cancer occurs in thousands of people, causing them great pain and distress, but it is rarely fatal. How best to maximize the good? On which disease should the government spend its time and treasure? Answering this question (and others like it) requires trying to apply the utility principle—a job often involving complex calculations of costs and benefits and frequently generating controversy.

Justice

In its broadest sense, *justice* refers to people getting what is fair or what is their due. In practice, most of us seem to have a rough idea of what justice entails in many situations, even if we cannot articulate exactly what it is. We know, for example, that it is unjust for a bus driver to make a woman sit in the back of the bus because of her religious beliefs, or for a judicial system to arbitrarily treat one group of citizens more harshly than others, or for a doctor to care for some patients but refuse to treat others just because he dislikes them.

Questions of justice arise in different spheres of human endeavor. *Retributive justice*, for example, concerns the fair meting out of punishment for wrongdoing. On this matter, some argue that justice is served only when people are punished for past wrongs, when they get their just deserts. Others insist that justice demands that people be punished not because they deserve punishment, but because the punishment will deter further unacceptable behavior. *Distributive justice* concerns the fair distribution of society's advantages and disadvantages—for example, jobs, income, welfare aid, health care, rights, taxes, and public service. Distributive justice is a major issue in bioethics, where many of the most intensely debated questions are about who gets health care, what or how much they should get, and who should pay for it.

Distributive justice is a vast topic, and many theories have been proposed to identify and justify the properties, or traits, of just distributions. A basic precept of most of these theories is what may plausibly be regarded as the core of the principle of justice: *Equals should be treated equally.* (Recall that this is one of the defining elements of ethics itself, impartiality.) The idea is that people should be treated the same unless there is a morally relevant reason for treating them differently. We would think it unjust for a physician or nurse to treat his White diabetic patients more carefully than he does his Black diabetic patients—and to do so without a sound medical reason. We would think it unfair to award the only available kidney to the transplant candidate who belongs to the "right" political party or has the best personal relationship with hospital administrators.

The principle of justice has been at the heart of debates about just distribution of benefits and burdens (including health care) for society as a whole. The disagreements have generally not been about the legitimacy of the principle, but about how it should be interpreted. Different theories of justice try to explain *in what respects* equals should be treated equally.

Libertarian theories emphasize personal freedoms and the right to pursue one's own social and economic well-being in a free market without interference from others. Ideally the role of government is limited to night-watchman functions—the protection of society and free economic systems from coercion and fraud. All other social or economic benefits are the responsibility of individuals. Government should not be in the business of

helping the socially or economically disadvantaged, for that would require violating people's liberty by taking resources from the haves to give to the have-nots. So universal health care is out of the question. For the libertarian, then, people have equal intrinsic worth, but this does not entitle them to an equal distribution of economic advantages. Individuals are entitled only to what they can acquire through their own hard work and ingenuity.

Egalitarian theories maintain that a just distribution is an *equal* distribution. Ideally, social benefits—whether jobs, food, health care, or something else—should be allotted so that everyone has an equal share. Treating people equally means making sure everyone has equal access to certain minimal goods and services. To achieve this level of equality, individual liberties will have to be restricted, measures that libertarians would never countenance. In a pure egalitarian society, universal health care would be guaranteed.

Between strict libertarian and egalitarian views of justice lie some theories that try to achieve a plausible fusion of both perspectives. With a nod toward libertarianism, these theories may exhibit a healthy respect for individual liberty and limit governmental interference in economic enterprises. But leaning toward egalitarianism, they may also mandate that the basic needs of the least well-off citizens be met.

In bioethics, the principle of justice and the theories used to explain it are constantly being marshaled to support or reject health care policies of all kinds. They are frequently used—along with other moral principles—to evaluate, design, and challenge a wide range of health care programs and strategies. They are, in other words, far from being merely academic.

ETHICAL RELATIVISM

The commonsense view of morality and moral standards is this: There are moral norms or principles that are valid or true for everyone. This claim is known as **moral objectivism**, the idea that at least some moral standards are objective. Moral objectivism, however, is distinct from **moral absolutism**, the belief that objective moral principles allow no exceptions or must be applied the same way in all cases and cultures. A moral objectivist can be absolutist about moral principles, or she can avoid absolutism by accepting that moral principles are prima facie. In any case, most people probably assume some form of moral objectivism and would not take seriously any claim implying that valid moral norms can be whatever we want them to be.

But moral objectivism is directly challenged by a doctrine that some find extremely appealing and that, if true, would undermine ethics itself: **ethical relativism**. According to this view, moral standards are not objective but are relative to what individuals or cultures believe. There simply are no *objective* moral truths, only *relative* ones. An action is morally right if endorsed by a person or culture and morally wrong if condemned by a person or culture. So euthanasia is right for person A if he approves of it but wrong for person B if she disapproves of it, and the same would go for cultures with similarly diverging views on the subject. In this way, moral norms are not discovered but made; the individual or culture makes right and wrong. Ethical relativism pertaining to individuals is known as **subjective relativism**, more precisely stated as the view that right actions are those sanctioned by a person. Ethical relativism regarding cultures is called **cultural relativism**, the view that right actions are those sanctioned by one's culture.

In some ways, subjective relativism is a comforting position. It relieves individuals of the burden of serious critical reasoning about morality. After all, determining right and wrong is a matter of inventorying one's beliefs, and any sincerely held beliefs will do. Morality is essentially a matter of personal taste, which is an extremely easy thing to establish. Determining what one's moral views are may indeed involve deliberation and analysis—but neither of these is a necessary requirement for the job.